

Do LLMs Silently Compute the Wrong Thing?

Measuring Silent Semantic Errors in LLM Data-Transform Code and Detecting Them with Specification-Derived Operator Contracts

Anonymous submission

Abstract

Large language models increasingly write the *data-transform* code behind analyses and figures—group-wise medians, weighted means, joins, pooled rates. We show that a large fraction of this code is *silently* wrong: it runs without error and produces plausibly-shaped output, yet computes a different quantity than the user asked for. The hazard is *executable and gold-free*: an under-specified request admits several defensible readings that *disagree materially* on the actual table, and the model *silently commits* to one without flagging the fork. On real Nature source-data tables, frontier models commit away from the natural reading in **77%** of *ambiguously*-phrased data-transform tasks (10% when disambiguated), and the failure is *invisible* to every simple safeguard: deterministic checks (execution, shape, range, self-consistency) catch **0%**, and LLM self-critique misses 39% while false-flagging 40%—a pervasive, invisible failure mode whose measurement is our primary contribution. As a goldless remedy where the operator is known or inferable, we introduce *specification-derived operator contracts*: executable checks over the (input, parameters, result) triple that certify an operator’s semantics using *no per-task gold answer*—the specification comes from the operator intent, not a stored label, and can be auto-synthesized from a natural-language goal. Across two real-data slices (up to 1,855 checks) they flag silent errors at **98–99% recall / 0.2% false-positive** under a *known* operator, where existing checks are structurally blind—about a third via *pure relational invariants* that compute no reference value (a *five-family* non-circular core; 99.8% on the large real slice) and the rest via *parameterized recomputation* that presupposes the operator (96.6%). We show the contracts can be synthesized from a high-level goal (78–83% correct, $N=23$, one-shot), transfer across execution substrates unchanged (pandas→SQL, 10/10 operators), and drive a zero-training repair recipe (75% vs. 12% for strong self-debug). Finally, we run the *entire* pipeline—infer operator, ground column roles, synthesize contract, detect—from only a messy request and a raw table, catching 51–58% of silent errors with nothing structured provided—a useful early-warning floor, far below the known-operator regime rather than a deployable replacement for it. We frame the work *measurement-first*—a common, invisible failure quantified, with operator intent (once recovered) making it checkable without a per-task gold answer—a *goldless executable veri-*

fier for the verifiable-reward and data-agent regime where no per-task reference exists—and release the measurement suite and code.

1 Introduction

Code-writing assistants now sit in the middle of everyday data analysis: a user describes what they want in plain language (“give me the typical expression per cell type”), and the model returns `pandas` (or `SQL`) that computes it. When that code has a *syntactic* bug it crashes; when it has an obvious numeric bug the plot looks wrong. The dangerous case is the third one: the code *runs*, the output has the right shape and a believable range, and it is nonetheless the answer to a *different question*—a mean where the user meant a median, a per-row ratio averaged instead of a pooled rate, an inner join that silently drops unmatched rows. We call these **silent semantic errors**.

Silent semantic errors are pernicious precisely because the standard safety net does not see them. Execution-based evaluation (“does it run?”), shape/type validity, and self-consistency across samples all pass, because the wrong code is still *valid* code producing *valid-looking* output. Verifying the *meaning* of the result normally requires a per-task gold reference—but in deployment there is none: if we already knew the right answer we would not be asking the model.

This paper. We make three moves. **(1) Measurement**—our primary contribution. We quantify how often frontier LLMs silently compute the wrong data transform, on *real* scientific data (Nature source-data tables) and across models, and quantify how completely existing checks miss it. **(2) A specification-derived detector** for the known- or inferred-operator regime. We introduce *typed operator contracts*: executable checks of an operator’s semantics—conservation laws, per-group identities, row-count and range constraints, or, where appropriate, an intent-derived recomputation—evaluated over the (input, params, result) triple with *no per-task gold label*. **(3) From assertions to an AI pipeline.** The contracts are not a hand-written wall of `asserts` but the core of a system that *grounds an informal request into a verifiable specification*: it infers the operator and column roles from a messy request, synthesizes the invariant from a natural-language goal, and repairs the code—end to end.

This is an anonymized submission for review purposes only. Distribution, citation, or public sharing of this manuscript is strictly prohibited. Copyright and publication details will appear in the final version if accepted.

Where this sits. The target is the reliability of *LLM data-analysis agents*—the systems now writing the `pandas/SQL` behind scientific figures and reports—and the work connects three current threads: ambiguity-aware code generation (Fang et al. 2026; Yang et al. 2026), verifiers and verifiable rewards, and reference-free (goldless) verification. Our angle is a goldless, executable, and *reward-hack-resistant* check for *consequential intent ambiguity* in real data analysis—a signal that can gate, flag, or (at inference time) reward.

What “no gold” means. A contract needs *no per-task human-labeled answer*—it derives from the operator’s *specification*. For some operators it is a *pure invariant* no implementation pins down (within-group shares sum to one *per group*; a left join keeps every left row); for others it *recomputes* the expected value from the known operator and checks agreement. The assumption is not “no reference is ever computed” but that the *operator and roles are known or inferred*—we measure separately what happens when they must be recovered (Sec. 7.3).

Contributions.

1. **A measurement study** of silent semantic errors in LLM data-transform code: 77% under author-designed ambiguous requests (10% disambiguated), 32–65% across eight models and four vendors, 26% on external DS-1000; deterministic checks detect 0%, self-critique 61% recall at 40% false-positive (Sec. 4).
2. **Specification-derived operator contracts** that, *given a known or inferred operator*, flag the errors at 98–99% recall / 0.2% false-positive on real tables—anchored by a *five-family* pure-invariant core that derives no reference value (99.8% on the large real slice), the rest presupposing the operator—where code-layer baselines are structurally blind (Sec. 6).
3. **Grounding intent into a specification:** the invariant can be auto-synthesized from a high-level goal (78–83%, $N=23$, one-shot; stable across GPT and Gemini), and the operator (98%) and column roles (74–83%) can be recovered from a keyword-free request (Sec. 7–7.3).
4. **An end-to-end pipeline** run on *real* Nature tables from only a messy request and a raw table—nothing structured given—catching 51–58% of silent errors, plus **substrate-agnostic transfer** (`pandas`→`SQL`, 10/10 operators, zero `SQL`-specific code) and a **zero-training repair recipe** (targeted 75% vs. 12% strong self-debug; Sec. 7.3–8).
5. **Honest, quantified boundaries:** the method’s power lives where operator-level intent is recoverable; on raw free-text code operator inference is the bottleneck, and a calibrated scope-gate makes out-of-scope inputs *safely abstain* (Sec. 9).

2 Related Work

Evaluating LLM code generation. HumanEval, MBPP, and DS-1000 (Chen et al. 2021; Austin et al. 2021; Lai et al. 2023) score functional correctness against *hidden tests* or a *gold reference*; execution-based agent evaluation (e.g., SWE-bench (Jimenez et al. 2024)) assumes a checkable oracle. Our setting is complementary and harder: no per-task reference

exists at inference time, and the code already passes execution.

Data validation and pipeline testing. Practitioner frameworks such as Great Expectations, `pandera`, and `dbt` tests (Great Expectations 2020; Bantilan 2020) let engineers assert schemas, ranges, and distributional expectations on tables. These are *hand-authored per dataset* and check *surface* properties (types, nullability, bounds)—exactly the properties that silent semantic errors preserve. We instead check *operator semantics* and *synthesize* the check from intent.

Property-based and metamorphic testing. Property-based testing (Claessen and Hughes 2000) and design-by-contract (Meyer 1992) predate us by decades; metamorphic testing (Chen, Cheung, and Yiu 1998) verifies relations between outputs without a gold oracle. Invariant mining (Daikon; Ernst et al. 2007) automatically discovers likely assertions from a corpus of *correct* executions; we instead derive the assertion from natural-language *intent*, with no correct execution to mine. Our mechanism is in this family. Our contribution is not the mechanism of asserting a property but (i) an *empirically-grounded* taxonomy of the operator-level slips LLMs actually make, (ii) checks that need *no per-task gold* and *can be synthesized from natural language*, and (iii) a pipeline that *grounds a free-text request into the right invariant*—turning property-based testing from a developer tool into an inference-time detector. Concretely, a generic PBT or metamorphic tool cannot flag a global-vs-per-group share slip on its own: *which* property to enforce (shares sum to one per group, not globally) *is* the user’s inferred intent. Selecting and instantiating that invariant from a messy request is the inference problem we solve; the assertion itself is deliberately old and simple.

Validating LLM outputs without gold. Self-consistency (Wang et al. 2023), self-critique / self-refine (Madaan et al. 2023), LLM-as-judge (Zheng et al. 2023), and execution-guided verification (e.g., LEVER Ni et al. 2023) estimate correctness without a reference. We show these miss silent semantic errors (self-critique: 61% recall at 40% false-positive); our consensus + metamorphic refinement is a goldless *verification signal* in the spirit of metamorphic testing.

Verifiers and verifiable rewards. Outcome verifiers (Cobbe et al. 2021), process reward models (Lightman et al. 2024), and RL from *verifiable* rewards (DeepSeek-AI et al. 2025) drive much of the recent gain in reasoning and code, and test-time verification turns extra inference compute into accuracy (Snell et al. 2024). Each assumes a checkable signal exists—unit tests, an answer key, or a learned reward. Because our contract is a *sound* executable invariant rather than a learned reward, it is far harder to *reward-hack* than model-based verifiers. We find it pays off at *test time* (goldless selection and repair feedback, Sec. 8) but, used as an offline preference label, does not improve an open model—optimizing a goldless (necessary-but-not-sufficient) signal Goodharts—a concrete data point on where goldless verifiers help and where they invite reward-hacking.

Approach	operator	semantic authoring no per-task	intent from NL	no gold label	infer.-time detector
Great Expectations / pandera	×	×	×	✓	✓
Property-based testing	~	×	×	~	×
Metamorphic testing	~	×	×	✓	×
LLM self-critique	✓	✓	✓	✓	✓
(but 61% recall / 40% FP)					
Ours	✓	✓	✓	✓	✓

Table 1: Positioning. Data-validation frameworks check *surface* properties (types/ranges) that silent errors preserve; PBT/metamorphic testing are developer-time and hand-authored; self-critique is inference-time but too noisy (Sec. 4). We check *operator semantics*, synthesized from intent, with no per-task label, as an inference-time detector. (~: partial/varies.)

LLMs for data science and visualization. Chart/plot benchmarks and code-/visual-similarity judges (MatPlotAgent, Plot2Code, ChartMimic, VisEval) (Yang et al. 2024; Wu et al. 2024; Shi et al. 2024; Chen et al. 2024) score whether a figure *looks* right, not whether the plotted numbers *equal* the data. We measure and detect the upstream data-transform error such judges cannot see. Table 1 summarizes these distinctions.

3 Problem Formulation

A *data-transform task* is a triple (D, ι, θ) : an input table (or tables) D , a natural-language intent ι , and column-role parameters θ (which column is the group, the value, the weight, ...). Real intents are routinely *under-specified*: “a typical value per group” admits a set of defensible operator readings $\mathcal{G} = \{g_1, \dots, g_k\}$ (mean vs. median; per-group vs. global share; pooled rate vs. mean-of-ratios; inner vs. left join). The task is **consequently ambiguous** on D when these readings *disagree materially*— $\exists g_i, g_j \in \mathcal{G} : \|g_i(D) - g_j(D)\| > \tau$ —an *executable, gold-free* property (run the readings and compare; on our operator fixtures all forks exceed $\tau=1\%$, often by orders of magnitude). A model emits code f with $r = f(D)$; the failure we target is a **silent commitment**: f executes, r is schema-valid, and the model commits to one reading *without surfacing the fork*. Whether the commitment is “wrong” need not appeal to a stored answer—the readings’ disagreement already establishes the hazard. As a *directional* refinement we take the natural-phrasing reading $g \in \mathcal{G}$ as a reference and call r a **silent semantic error** when it commits away from it ($r \neq g(D)$); g is a reference for that directional rate, *not* a claim that it is the uniquely correct answer.

A **specification-derived contract** for an operator o is a function $C_o(D, \theta, r) \in \{\text{fire}, \text{pass}\}$ that identifies *which* reading r computed and fires when it is a confusable alterna-

tive to o ’s intended semantics. C_o uses *no per-task gold label*: it is a function of the operator specification alone (instantiated at the given θ). We evaluate it on held-out correct and slip implementations by two properties: it *fires on the silent slip* (recall) and *passes on every valid implementation*, including order/representation-diverse correct ones (false-positive robustness). Crucially C_o reads only (D, θ, r) —never the code f —so it is agnostic to how r was produced (Sec. 7.4). The strong assumption is that the operator o and roles θ (the reading menu) are known or inferable; Sec. 7.3 removes it by inferring them from ι .

4 Measurement: How Common, How Invisible

Corpus and protocol. We pair Nature figures with their Source-Data tables and instantiate operator-semantic tasks (median-vs-mean, within-group share, weighted mean, pooled rate, NaN-as-zero, ...) on the real tables, so the tables and their columns are real scientific data, *not* authored by us; the operator set (24 operators grounded in observed real-data slips, not chosen post hoc; released), its gold transform, and the ambiguous/clarified prompts are fixed author-written templates instantiated on those tables (table-level independence, not task-level). Each task carries a matched pair of prompts: an *ambiguous* phrasing (“a typical value per group”—the natural way a hurried analyst asks) and a *clarified* phrasing that names the disambiguation. Gold is the template transform; a response is scored *silent* if it runs, is schema-valid, and disagrees with gold. We report over a 71-article independent sample (each article contributes a capped number of tasks) with 95% Wilson intervals.

Prevalence. Under the ambiguous phrasing, frontier models emit a silent semantic error on **77%** of real Nature tasks (1296/1682, CI 75–79); disambiguating the request drops this to **10%** (175/1682, CI 9–12; 21% on a larger 229-article slice). The gap is the point: much of the danger is *under-specification* that the model resolves the wrong way, and the rate is **40–60%** across five frontier models and three vendors on the expansion grid (Table 3; **32–42%** for the same models on the core grid), rises to **65%** for smaller open models (Qwen-2.5-Coder-7B, falling to 44% at 32B), and is **26%** on external DS-1000, so it is a property of the task family, not one model (Fig. 1).

Invisibility. We measure what existing safeguards catch on the operator grid (Table 2). Deterministic checks—execution, shape/type, value-range, multi-sample consistency—catch **0%**: the wrong code runs, is shaped correctly, and is *stably* wrong. Multi-sample consistency stays at 0% even when the K samples are drawn at temperature 0.8 (recall 0/38): the silent errors are *common-mode*—the model resamples the same wrong transform (92% of tasks give identical output even at $T=0.8$), so disagreement-based checks never fire. LLM self-critique—the same model, shown the *clarified* request and a preview of the produced table and asked in one turn for a JSON correct/incorrect verdict—catches 61% but false-flags 40%, unusable as a gate; note this is a *favorable* setup (the judge is handed the disambiguated intent, not the




Figure 1: Silent-error *prevalence* (red: ambiguous 77% vs. clarified 10%) and *detection recall* (blue baselines vs. our contracts, green). Deterministic checks detect 0%; self-critique’s 61% comes with 40% false-positives.

Safeguard	Recall	False-positive
Execution/shape/range/consistency	0%	0%
LLM self-critique	61%	40%
Specification contracts (ours)	98–99%	0.2%

Table 2: Detection of silent semantic errors. Deterministic checks and self-critique are shown on the operator grid; a 119-task real-Nature re-run gives the same profile (0% recall / 62% recall at 67% FP), so the 0% is not a synthetic artifact. The **ours** row (98–99%) is on real Nature data (Sec. 6). Deterministic checks are structurally blind; self-critique is too noisy to gate on.

ambiguous one) and it still fails. Re-running the same five detectors *directly* on the real Nature slice (119 tasks) confirms this is not an artifact of the synthetic grid: execution, validity, and multi-sample consistency stay at 0% recall on real data, and self-critique is if anything worse in the wild (62% recall at 67% false-positive), while our contracts hold at 100%/0%.

5 Specification Contracts as an AI Pipeline

5.1 Typed operator contracts

Each operator is certified by one small check (≈ 10 lines, reading only (D, θ, r)). Examples: a *weighted mean* must equal $\sum v_i w_i / \sum w_i$; *within-group shares* must sum to one *per group* (not globally); a *pooled rate* equals $\sum \text{num} / \sum \text{den}$, not the mean of per-row ratios; a *left join* preserves every left row. Contracts range from pure invariants (share, join) to intent-derived recomputations (weighted mean); none uses a per-task human label. The taxonomy of operators is *empirically grounded*—derived from the slip patterns the measurement study documents—not an arbitrary catalog, and (below) it is auto-extensible.

5.2 Synthesizing the specification from language

To show the taxonomy is not a closed hand-written set, we ask a frontier model to *synthesize* the contract from only the NL goal, the parameter names, and the result schema—never

Model	Vendor	Silent rate	Our FP
GPT-5.5	OpenAI	60%	0/24
GPT-5.4	OpenAI	55%	0/21
GPT-5.3-codex	OpenAI	55%	0/23
Claude-Opus-4.8	Anthropic	55%	0/23
Gemini-3.1-Pro	Google	40%	0/28

Table 3: Silent-error rate is a property of the task family, not one model: 40–60% across five frontier models from three vendors on the harder expansion-operator grid ($N=20/\text{model}$; the same models score 32–42% on the core grid). Per-model rates ($N=20\text{--}23$) are directional; primary claims pool to $N=115$ (Sec. 7.3). The contract’s false-positive rate is **0** on *every* vendor—the detector is not tuned to one model’s output style.

the reference implementation. To rule out the model merely *transcribing a stated formula*, we run a *de-leaked* condition whose intent gives only a high-level goal—no formula and no “not-the- X ” slip-contrast (a plain-language goal may still *imply* the target property for some operators, a residual we audit in Sec. 7). We score each synthesized contract **CORE** (fires on the silent slip *and* passes on the correct implementation) and **FULL** (CORE *and* robust to alternative valid implementations).

5.3 Goldless counterexample-guided refinement

Synthesis can over-fire on valid alternatives. We add a monotone-safe refinement loop with a goldless oracle (Fig. 2): sample K diverse implementations of the intent, cluster by output into a *consensus* (probably-correct) set, and treat any contract that fires on a consensus member as a false-positive counterexample; *metamorphic* perturbations of a consensus result probe non-triviality. A refinement is accepted only if a goldless score (consensus-pass + probe-fire) strictly improves, so the loop never regresses below one-shot by construction. We report this as a *mechanism with a safety guarantee*, not an accuracy gain (Sec. 7).

5.4 From a messy request to a check

Deployed, the system receives only a messy request ι and a raw table D (Fig. 3). It (1) infers the operator $\hat{o}$ from ι , (2) grounds the column roles $\hat{\theta}$ from ι and the column names, (3) synthesizes $C_{\hat{o}}$, and (4) checks r . When a contract fires, the same signal drives *repair*: the violated invariant is fed back as a targeted prompt, and best-of- N selection uses the contract (not gold) to pick among candidate fixes. Figure 4 walks through one real case end to end.

6 Detection on Real Data

Unless noted, the frontier model is a 2026-era GPT model behind an OpenAI-compatible API; buggy starting points for repair are *model-generated*, not author-written; intervals are 95% Wilson.¹

¹Model identifiers (“GPT-5.4/5.5”, “Gemini-3.1-Pro”) are the API names available at submission time; we release prompts, seeds, and all outputs for reproducibility.

Goldless refinement of a synthesized contract
input intent ι , params θ , input D ; **output** contract C

```

1  $C \leftarrow \text{SYNTHESIZE}(\iota, \theta)$  // one-shot
2  $\{f_k\} \leftarrow \text{DIVERSEIMPLS}(\iota, \theta, K)$ 
3  $S \leftarrow$  largest output-agreement cluster of  $\{f_k(D)\}$  // consensus
4  $P \leftarrow \text{PERTURB}(s)$  for  $s \in S$  // metamorphic probes
5 repeat
6   if  $C$  fires on some  $s \in S$  then // FP counterexample
7      $C' \leftarrow \text{RESYNTHESIZE}(\iota, \theta, s)$ 
8     if  $\text{score}(C'; S, P) > \text{score}(C; S, P)$  then  $C \leftarrow C'$ 
9 until no accepted change return  $C$ 
 $\text{score}(C; S, P) = \#\{s \in S : C \text{ passes}\} + \#\{p \in P : C \text{ fires}\}$ ; no gold used.

```

Figure 2: Monotone-safe goldless refinement. The consensus set S and metamorphic probes P replace a reference answer; line 8 guarantees the loop never regresses below the one-shot contract *on the goldless score* (not necessarily on downstream detection quality).




Figure 3: End-to-end pipeline. From only a messy request and a raw table, the system infers the operator (98%), grounds column roles (74–83%), synthesizes the contract (78%), and checks the result—nothing structured is provided.

The contracts detect silent errors at **98% recall** (1438/1471, CI 97–98) and **0.2% false-positive** (4/1855, CI 0.1–0.6) on a 71-article independent sample; a second, table-dense slice gives 1398/1408 recall (99%) and 2/1539 false-positive. We report the more conservative 71-article number as the headline because it maximizes cross-paper independence. We stress a distinction the framing depends on: this measures detection under *known* operator and parameters; the self-authored correctness grid used in development is a contract *unit test*, not evidence of generalization, and we exclude it from headline claims. We report the false-positive rate as the exact fraction 4/1855 (0.2%) with its interval throughout, rather than rounding to “0%”.

Failure analysis. The $\sim 2\%$ of silent errors the contracts miss are not contract bugs but an inherent *detectability* limit: they concentrate where the slip’s output is numerically near-identical to the correct one—e.g., a group whose values are near-uniform (so mean $\approx$ median) or a weight column that is near-constant (so weighted $\approx$ plain mean). There, *no* invariant can separate the two within tolerance, and the “error” has negligible practical effect. The four false-positives arise on degenerate tables (a single group, or all-equal values) where the invariant is *undefined* on the correct output (e.g. a per-group share of $0/0=\text{NaN}$), so the contract fires on a valid result; a schema/degeneracy pre-filter removes them.

Request: “add a column with each row’s sales as its part of the total sales for that region”

```

(1) infer operator  $\rightarrow$  within_group_share
(2) ground roles  $\rightarrow$  group=region, value=sales
(3) synthesize contract (from the goal, no gold):
    per = result.groupby(g)[share].sum()
    fire if any(|per - l| > tol) # per-GROUP
(4) model’s code (silent slip):
    df[share] = df[sales] / df[sales].sum() # GLOBAL
(5) contract FIRES: “shares sum to 1 globally,
    should be per-group”
(6) targeted repair  $\rightarrow$ 
    df[share] = df[sales] /
    df.groupby(region)[sales].transform(sum) # PASS

```

Figure 4: A worked example on a real Nature table. The same specification catches a `within_group_share` silent slip (global instead of per-group total) and drives the targeted repair—no per-task gold, operator and roles inferred from the request. This is a semantic invariant, not a hand-written per-table assertion.

Operator	Recall on real silent errors
weighted-mean	140/140 = 100%
nan-as-zero-sum	16/16 = 100%
within-group-share	538/539 = 99.8%
pooled-rate	118/119 = 99%
median-not-mean	626/657 = 95%

Table 4: Per-operator detection recall on the real Nature sample. High and uniform; median-not-mean’s 95% is the near-tie detectability limit (Failure analysis), not a contract defect.

Per-operator breakdown. Detection is uniform across operators, not driven by one easy case (Table 4): recall is $\geq 99\%$ for four of five operators and is only pulled down by median-not-mean (95%), exactly the near-tie regime the failure analysis predicts (groups where the median and mean nearly coincide). Which operator carries the residual misses is *slice-dependent*—on other real slices median-not-mean is 100% and the misses fall on within-group-share/pooled-rate—so the invariant claim is that residual misses are near-ties at the chosen tolerance, not that median is intrinsically the hardest operator.

Two kinds of contract (what “no gold” really buys). We separate recall by contract class, because the two make different epistemic claims and a reviewer should see the split rather than a pooled number. A *pure invariant* certifies a relational property that no single implementation pins down and derives *no expected value*—within-group shares must sum to one *per group*, a left join must preserve every left row, a running total’s last entry must equal the column sum. A *recomputation* instead derives the expected value from the (known or inferred) operator and parameters and checks agreement—a weighted mean equals $\sum v_i w_i / \sum w_i$, a pooled rate equals $\sum \text{num} / \sum \text{den}$. We concede the latter amounts to a *parameterized reference computed on the fly*, not merely an invari-

Intent given to synthesizer	CORE	FULL
Formula-stating (“leaky”)	83% [63–93]	78% [58–90]
High-level goal only (de-leaked)	78% [58–90]	70% [49–84]

Table 5: One-shot goldless synthesis, $N=23$ operators. **CORE**: fires on the slip and passes on the correct impl. **FULL**: CORE and robust to alternative valid impls. Intervals overlap at $N=23$; we claim feasibility, not a significant de-leak gap.

ant: its power is real but conditional on knowing the operator. On the real Nature sample the *pure-invariant* contract (within-group-share) detects at $538/539 = 99.8\%$. Crucially this non-circular core is *not one operator*: five core contracts derive no reference value—within-group -share, left-join-keep-all, running-total-equals-column-sum, top- n -with-ties, and count-includes-empty—and on the shared fixtures each fires on its silent slip while passing the intended transform *and* every order/representation-diverse valid implementation (recall 5/5, false-positive 0; within-group-share is simply the family member with a large real-data sample). The *recomputation* contracts (weighted-mean, pooled-rate, median, nan-as-zero) detect at $900/932 = 96.6\%$. Both need *no per-task human-labeled answer*; the honest distinction is whether a reference *value* is derived, and we state plainly that roughly two-thirds of the headline checks are recomputations whose strength presupposes a known operator (Sec. 7.3 measures what happens when it must be inferred).

7 Synthesis, Grounding, and End-to-End

7.1 Synthesizing contracts from a goal

Table 5 reports one-shot synthesis over 23 operators. Removing the formula from the intent costs only 5 CORE points ($83 \rightarrow 78$); we are careful *not* to over-read this small- N gap (the Wilson intervals overlap), and claim only that goldless synthesis from a high-level goal is *feasible at useful rates*, not that de-leaking is free. The result is stable across GPT-5.4, GPT-5.5, and (different vendor) Gemini-3.1-Pro at $\approx 83\%$ CORE, which argues against model cherry-picking. The goldless refinement loop is monotone-safe and *neutral* on this already-strong one-shot synthesizer (both 83% CORE, $N=12$); we include it as a verification mechanism with a no-regression guarantee, not as an accuracy contribution.

7.2 Recovering operator and roles from language

Given a realistic request with *no operator keyword*, the frontier model recovers the operator on real Nature tasks at **98%** (147/150, CI 94–99), while a keyword regex collapses to **21%** (31/150)—the model uses semantics, not string matching (the keyword-stuffed upper bound is 100% for both). Column-role grounding from the request reaches **83%** on type-disambiguated roles (categorical group vs. numeric value) and 74% overall; the shortfall is genuinely symmetric roles (a weighted mean’s value-vs-weight, a pooled rate’s numerator-vs-denominator) that the request does not uniquely fix—an *ambiguity*, not a method failure.

Setting	Op infer	Contract	Full system
Synthetic fixtures ($N=23$)	83%	74%	61% [41–78]
Real Nature tables ($N=33$)	76%	70%	58% [41–73]
+ inferred roles ($N=35$)		roles 74/83%	51% [36–67]
Cross-vendor pooled ($N=115$, 5 vendors)	80%	70%	57% [48–66]

Table 6: End-to-end zero-template detection. Operator and contract are inferred throughout; the $N=35$ row additionally infers column roles (only the messy request and raw table are given). The cross-vendor pooled row (five models across OpenAI, Google, and Anthropic, $N=115$; column roles given) tightens the interval; per-vendor full-system spans 48–70%.

7.3 End-to-end from only a request and a table

Table 6 runs the full pipeline with nothing template-given. On generated fixtures it catches 61%; on *real* Nature tables 58%; and with column roles also inferred (only a messy sentence and a raw table), 51%—3 points below the params-given setting. **We treat 51–58% as a boundary, not a headline**: roughly half of silent errors still pass, so the deployable regime remains the structured-intent 98%/0.2%; the value here is scientific—the pipeline degrades gracefully and localizes where accuracy is lost (see below).

Where the loss is. Pooling across five vendors ($N=115$) tightens the full-system estimate to **57%** [48–66] (width 18 vs. 38 at $N=23$) and lets us *decompose* the miss: of 49 misses, **53%** are contract-synthesis failures, **31%** operator-inference, 16% both—so invariant synthesis, not operator inference, is the dominant gap (26 vs. 15; intervals overlap, directional) and the concrete future-work target. De-leaked synthesis alone at $N=115$ holds at 80% [72–86] CORE (GPT-5.5 96%, codex 87%, GPT-5.4 78%, Gemini 74%, Opus 65%), so the feasibility number is not a single-model artifact.

Exemplar prompting does not fix synthesis (ablation). Showing the synthesizer a worked contract does *not* help: over $k=3$ runs/vendor, zero-shot is the most stable arm (pooled 72%; GPT $83 \pm 0\%$) while one/few-shot fall to $57 \pm 9\%/54 \pm 7\%$ —within run-to-run noise, so zero-shot is the right default. Residual failures concentrate on mechanically-intricate operators (index alignment, dtype coercion, resample boundaries)—a fine-tuning/tool-use target, not a prompting fix.

7.4 Substrate-agnostic transfer (pandas $\rightarrow$ SQL)

Because a contract reads only (D, θ, r) , the *same* contracts fire on the identical semantic slip written in SQL. Over 10 operators (weighted AVG, WHERE g IS NOT NULL, join fan-out, inner-join drop, AVG of ratios, missing DISTINCT, COUNT(`col`) vs COUNT(`*`), window OVER() vs PARTITION BY, ...) the contracts fire on every SQL slip and pass every correct SQL, with **zero** SQL-specific code (10/10, CI 72–100). We do not over-claim: these invariants are conservation/range/row-count properties that

Repair arm ($N=65$, 4 models)	Fixed	Over-rep.
generic (retry only)	18% [11–30]	0%
strong self-debug	12% [6–22]	0%
localization-only	52% [40–64]	0%
targeted (contract invariant)	75% [64–84]	3%
ceiling (gold-diff oracle)	92% [83–97]	6%

Table 7: Contract-guided repair (buggy starts are model-generated; one run over four models, all arms on the same $N=65$ silent starts). The invariant lifts repair from 12% (strong self-debug) to 75% with disjoint CIs and 3% over-repair; self-debug falling *below* generic retry (12 vs. 18%) shows an external specification is necessary, and the +23-point gap over localization-only shows the invariant *content* matters, not just location.

transfer by construction; operators whose invariant relies on pandas-specific NaN semantics would need restating.

8 Zero-Training Repair

Feeding the violated invariant back as a targeted repair prompt fixes silent errors far better than strong goldless self-repair (Table 7). Targeted repair reaches **75%** vs. **12%** for a strong self-debug baseline (per-value re-derivation + decision enumeration + self-critique), with *disjoint* intervals; localization-only reaches 52%, so the invariant *content*—not merely knowing *where* to look—adds a further +23 points. A telling negative result: strong self-debug (12%) underperforms even generic retry (18%), because it confidently “re-derives” the same wrong operator; only an *external* specification breaks the loop. Best-of- N selection using the contract (no gold) adds further gains with zero mis-repairs.

The contract as a verifier: test time vs. training. We probe the contract as a reusable goldless *verifier signal* on open models, where retraining is feasible (Qwen-2.5-Coder; held-out silent starts, gold hidden). Using the contract only to *select* among N goldless repair samples lifts the fix rate from 28% to **48%** (best-of-6, contract-pass 0.72, zero over-repair), and the gain *grows with scale* (29% at 3B $\rightarrow$ 48% at 7B)—a goldless instance of test-time verification (Snell et al. 2024). Yet using the *same* contract to label preference pairs for offline DPO does *not* help—it moves the 7B by -4 to -8 points. The goldless verifier therefore pays off at inference time (selection, repair feedback), not yet as a training signal—a targeted, honest data point for the verifiable-reward program (Sec. 2).

9 Limitations and Honest Boundaries

Conditional validity. The method’s strength lives where operator-level intent is recoverable. On *raw free-text code* (DS-1000), inferring which operator to check is the bottleneck and no operating point reaches usable recall/precision; a calibrated scope-gate instead makes out-of-scope inputs *abstain*, cutting spurious firing from 55% to 5%. The deployable claim is “NL intent recoverable,” not “arbitrary code.” **Taxonomy coverage.** Contracts detect the slips they specify; we report coverage honestly and treat the self-authored

grid as a unit test, not evidence of generalization—the real-data 98%/0.2% shows detection generalizes across real input *tables* under a known operator (the pure-invariant contract is the non-circular core), while *cross-operator* generalization is Sec. 7.3. **Small samples.** Per-condition slices start small ($N=23-35$); we mitigate by pooling across vendors—synthesis and end-to-end to $N=115$ (five vendors), repair to $N=65$ (four models) (Sec. 7.3, interval width $38 \rightarrow 18$)—but still report every CI rather than significance. **Symmetric roles.** For operators with interchangeable numeric roles the request may not fix which column is which; we score this as ambiguity. **End-to-end ceiling.** The 51–58% zero-structure figure is an early-warning floor, well below the structured-intent regime, and its dominant failure is invariant synthesis, not operator inference. **Ecological validity of the prevalence measurement.** The 77% headline is the paper’s primary contribution, so what it is conditioned on matters: the Nature *tables* are real, but the ambiguous *requests* are author-written templates designed to model how a hurried analyst under-specifies intent, not a sampled corpus of deployed analyst prompts. The clarified-phrasing drop to 10% shows the ambiguity—not task difficulty—drives the errors, but it does not establish how often real requests are this ambiguous; quantifying the ambiguity distribution of deployed requests is the key threat to external validity. A pilot on 40 real StackOverflow-derived requests (DS-1000, ambiguity-prone families, each reduced to its prose intent) indicates operator ambiguity is common in the wild—flagged for 22% [12–38] by a conservative keyword detector and 80% [65–90] by an LLM judge (single free model; per-item labels released)—so the ambiguous regime is ecologically common rather than a contrived trap, though the precise deployment distribution remains future work. **Gold-label reliability.** Gold is the author-written template transform; we did not run a second-rater agreement study, so the 4/1855 false-positives—attributed above to degenerate tables—are not independently adjudicated. **Reproducibility and precision.** Frontier generations use default API sampling (reasoning models are not bit-reproducible even at temperature 0), so we repeat the synthesis/ablation $k=3\times$ and report mean $\pm$ sd: zero-shot synthesis is stable (GPT 83 $\pm$ 0%) while the exemplar effect is within run-to-run noise (Sec. 7). End-to-end and repair remain single draws whose stochasticity exceeds the task-clustered Wilson intervals; fuller averaging is future work.

10 Conclusion

LLMs silently compute the wrong data transform far more often than execution- or shape-based checks reveal (77% under ambiguous requests; 0% caught by simple safeguards). Specification-derived operator contracts detect these errors without a per-task gold reference (98–99%/0.2%), can be synthesized from a high-level request, transfer across execution substrates, and drive a zero-training repair recipe—and the whole pipeline runs end-to-end from only a messy request and a raw table. We frame silent-error detection as *grounding an informal intent into a verifiable specification*, and release the measurement suite and code.

References

- Austin, J.; Odena, A.; Nye, M.; et al. 2021. Program Synthesis with Large Language Models. *arXiv:2108.07732*.
- Bantilan, N. 2020. pandera: Statistical Data Validation of Pandas Dataframes. In *SciPy*.
- Chen, M.; Tworek, J.; Jun, H.; et al. 2021. Evaluating Large Language Models Trained on Code. *arXiv:2107.03374*.
- Chen, N.; Zhang, Y.; Xu, J.; et al. 2024. VisEval: A Benchmark for Data Visualization in the Era of Large Language Models. *IEEE TVCG*.
- Chen, T. Y.; Cheung, S. C.; and Yiu, S. M. 1998. Metamorphic Testing: A New Approach for Generating Next Test Cases. *Technical Report HKUST-CS98-01*.
- Claessen, K.; and Hughes, J. 2000. QuickCheck: A Lightweight Tool for Random Testing of Haskell Programs. In *ICFP*.
- Cobbe, K.; Kosaraju, V.; Bavarian, M.; et al. 2021. Training Verifiers to Solve Math Word Problems. *arXiv:2110.14168*.
- DeepSeek-AI; Guo, D.; Yang, D.; et al. 2025. DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning. *arXiv:2501.12948*.
- Ernst, M. D.; Perkins, J. H.; Guo, P. J.; McCamant, S.; Pacheco, C.; Tschantz, M. S.; and Xiao, C. 2007. The Daikon System for Dynamic Detection of Likely Invariants. *Science of Computer Programming*, 69(1–3): 35–45.
- Fang, Z.; Jin, D.; Dong, Y.; Li, Y.; Zhang, K.; Jin, Z.; and Li, G. 2026. ClarifyCodeBench: Evaluating LLMs on Clarifying Ambiguous Requirements for Code Generation. *arXiv:2607.00711*.
- Great Expectations. 2020. Great Expectations: Always Know What to Expect From Your Data. <https://greatexpectations.io>.
- Jimenez, C. E.; Yang, J.; Wettig, A.; et al. 2024. SWE-bench: Can Language Models Resolve Real-World GitHub Issues? In *ICLR*.
- Lai, Y.; Li, C.; Wang, Y.; et al. 2023. DS-1000: A Natural and Reliable Benchmark for Data Science Code Generation. In *ICML*.
- Lightman, H.; Kosaraju, V.; Burda, Y.; et al. 2024. Let’s Verify Step by Step. In *ICLR*.
- Madaan, A.; Tandon, N.; Gupta, P.; et al. 2023. Self-Refine: Iterative Refinement with Self-Feedback. In *NeurIPS*.
- Meyer, B. 1992. Applying Design by Contract. *Computer*, 25(10): 40–51.
- Ni, A.; Iyer, S.; Radev, D.; et al. 2023. LEVER: Learning to Verify Language-to-Code Generation with Execution. In *ICML*.
- Shi, C.; Yang, C.; Liu, Y.; et al. 2024. ChartMimic: Evaluating LMM’s Cross-Modal Reasoning Capability via Chart-to-Code Generation. *arXiv:2406.09961*.
- Snell, C.; Lee, J.; Xu, K.; and Kumar, A. 2024. Scaling LLM Test-Time Compute Optimally Can Be More Effective Than Scaling Model Parameters. *arXiv:2408.03314*.
- Wang, X.; Wei, J.; Schuurmans, D.; et al. 2023. Self-Consistency Improves Chain of Thought Reasoning in Language Models. In *ICLR*.
- Wu, C.; Ge, Y.; Guo, Q.; et al. 2024. Plot2Code: A Comprehensive Benchmark for Evaluating Multi-modal LLMs in Code Generation from Scientific Plots. *arXiv:2405.07990*.
- Yang, D.; Xie, X.; Yang, X.; Hu, M.; Huang, Y.; Zhang, Y.; Miao, W.; Su, T.; Wan, C.; and Pu, G. 2026. Assessing the Impact of Requirement Ambiguity on LLM-based Function-Level Code Generation. *arXiv:2604.21505*.
- Yang, Z.; Zhou, Z.; Wang, S.; et al. 2024. MatPlotAgent: Method and Evaluation for LLM-Based Agentic Scientific Data Visualization. In *ACL Findings*.
- Zheng, L.; Chiang, W.-L.; Sheng, Y.; et al. 2023. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. In *NeurIPS*.

Reproducibility Checklist

1. General Paper Structure

- 1.1. Includes a conceptual outline and/or pseudocode description of AI methods introduced (yes/partial/no/NA) **yes**
- 1.2. Clearly delineates statements that are opinions, hypothesis, and speculation from objective facts and results (yes/no) **yes**
- 1.3. Provides well-marked pedagogical references for less-familiar readers to gain background necessary to replicate the paper (yes/no) **yes**

2. Theoretical Contributions

- 2.1. Does this paper make theoretical contributions? (yes/no) **no**
If yes, please address the following points:
 - 2.2. All assumptions and restrictions are stated clearly and formally (yes/partial/no) **NA**
 - 2.3. All novel claims are stated formally (e.g., in theorem statements) (yes/partial/no) **NA**
 - 2.4. Proofs of all novel claims are included (yes/partial/no) **NA**
 - 2.5. Proof sketches or intuitions are given for complex and/or novel results (yes/partial/no) **NA**
 - 2.6. Appropriate citations to theoretical tools used are given (yes/partial/no) **NA**
 - 2.7. All theoretical claims are demonstrated empirically to hold (yes/partial/no/NA) **NA**
 - 2.8. All experimental code used to eliminate or disprove claims is included (yes/no/NA) **NA**

3. Dataset Usage

- 3.1. Does this paper rely on one or more datasets? (yes/no) **yes**
If yes, please address the following points:
 - 3.2. A motivation is given for why the experiments are conducted on the selected datasets (yes/partial/no/NA) **yes**
 - 3.3. All novel datasets introduced in this paper are included in a data appendix (yes/partial/no/NA) **partial**

- 3.4. All novel datasets introduced in this paper will be made publicly available upon publication of the paper with a license that allows free usage for research purposes (yes/partial/no/NA) [yes](#)
- 3.5. All datasets drawn from the existing literature (potentially including authors' own previously published work) are accompanied by appropriate citations (yes/no/NA) [yes](#)
- 3.6. All datasets drawn from the existing literature (potentially including authors' own previously published work) are publicly available (yes/partial/no/NA) [yes](#)
- 3.7. All datasets that are not publicly available are described in detail, with explanation why publicly available alternatives are not scientifically satisfying (yes/partial/no/NA) [NA](#)

- 4.13. This paper lists all final (hyper-)parameters used for each model/algorithm in the paper's experiments (yes/partial/no/NA) [partial](#)

4. Computational Experiments

- 4.1. Does this paper include computational experiments? (yes/no) [yes](#)

If yes, please address the following points:

- 4.2. This paper states the number and range of values tried per (hyper-) parameter during development of the paper, along with the criterion used for selecting the final parameter setting (yes/partial/no/NA) [partial](#)
- 4.3. Any code required for pre-processing data is included in the appendix (yes/partial/no) [yes](#)
- 4.4. All source code required for conducting and analyzing the experiments is included in a code appendix (yes/partial/no) [yes](#)
- 4.5. All source code required for conducting and analyzing the experiments will be made publicly available upon publication of the paper with a license that allows free usage for research purposes (yes/partial/no) [yes](#)
- 4.6. All source code implementing new methods have comments detailing the implementation, with references to the paper where each step comes from (yes/partial/no) [partial](#)
- 4.7. If an algorithm depends on randomness, then the method used for setting seeds is described in a way sufficient to allow replication of results (yes/partial/no/NA) [yes](#)
- 4.8. This paper specifies the computing infrastructure used for running experiments (hardware and software), including GPU/CPU models; amount of memory; operating system; names and versions of relevant software libraries and frameworks (yes/partial/no) [partial](#)
- 4.9. This paper formally describes evaluation metrics used and explains the motivation for choosing these metrics (yes/partial/no) [yes](#)
- 4.10. This paper states the number of algorithm runs used to compute each reported result (yes/no) [yes](#)
- 4.11. Analysis of experiments goes beyond single-dimensional summaries of performance (e.g., average; median) to include measures of variation, confidence, or other distributional information (yes/no) [yes](#)
- 4.12. The significance of any improvement or decrease in performance is judged using appropriate statistical tests (e.g., Wilcoxon signed-rank) (yes/partial/no) [partial](#)